

Prediction of a Chemical whether it is musk or non musk, using Artificial Neural Network

Resources:

- Dataset: <https://drive.google.com/file/d/1pZhzZnaPi74aKCQImSPzrTxWzVeE0qv/view?usp=sharing>
- Anaconda navigator 2019.10 :<https://www.anaconda.com/distribution/> 3.7 :
- Spyder: <https://www.spyder-ide.org>
- pandas: <https://pandas.pydata.org/pandas-docs/stable/>
- Numpy: <https://docs.scipy.org/doc/numpy-dev/user/quickstart.html>
- Scikit-learn: <http://scikit-learn.org/stable/>
- Matplotlib: <https://matplotlib.org/2.1.0/index.html>
- Keras: <https://keras.io/>

Tools Used:

- Data Analysis and Preprocessing:
 - Pandas, Numpy and Scikit-learn
- Data Visualization:
 - Matplotlib
- Deep Learning Library:
 - Keras (with Tensorflow backend)

We will divide the project into 2 parts:

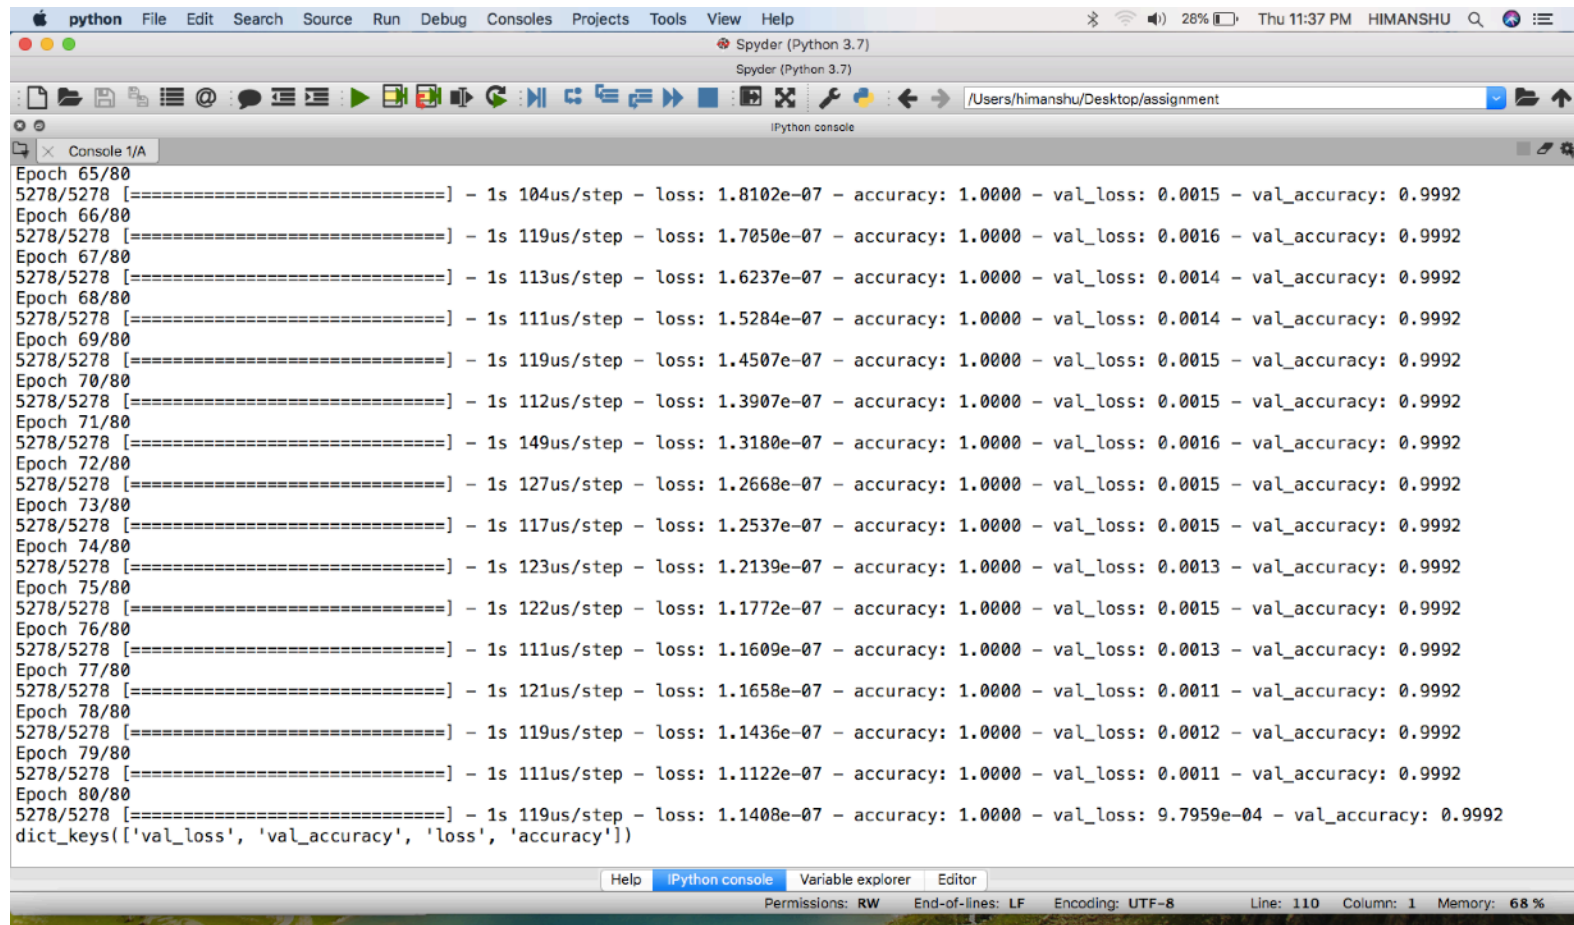
1. Data Exploration and Data Preprocessing:
 1. Description of Data
 2. Feature scaling with class StandardScaler
 3. Splitting data into training and test set
 4. Using matplotlib to visualization

Approach:

Network Architecture:

- Hidden layer with 84 Neurons(average of input and output nodes)
- Input layer with 166 neurons(number of independent features)
- First Hidden Layer with 84 Neurons and “Rectified” Activation Function
- Second Hidden Layer with 84 Neurons and “Rectified” Activation Function
- Output Layer with 1 Neurons and “Sigmoid” Activation Function
- Number of epochs 80 and batch size 20

Run time :



The screenshot shows the Spyder Python IDE interface. The top menu bar includes File, Edit, Search, Source, Run, Debug, Consoles, Projects, Tools, View, and Help. The toolbar contains various icons for file operations, running, and debugging. The main window displays a Jupyter console with the following output:

```
Epoch 65/80
5278/5278 [=====] - 1s 104us/step - loss: 1.8102e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 66/80
5278/5278 [=====] - 1s 119us/step - loss: 1.7050e-07 - accuracy: 1.0000 - val_loss: 0.0016 - val_accuracy: 0.9992
Epoch 67/80
5278/5278 [=====] - 1s 113us/step - loss: 1.6237e-07 - accuracy: 1.0000 - val_loss: 0.0014 - val_accuracy: 0.9992
Epoch 68/80
5278/5278 [=====] - 1s 111us/step - loss: 1.5284e-07 - accuracy: 1.0000 - val_loss: 0.0014 - val_accuracy: 0.9992
Epoch 69/80
5278/5278 [=====] - 1s 119us/step - loss: 1.4507e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 70/80
5278/5278 [=====] - 1s 112us/step - loss: 1.3907e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 71/80
5278/5278 [=====] - 1s 149us/step - loss: 1.3180e-07 - accuracy: 1.0000 - val_loss: 0.0016 - val_accuracy: 0.9992
Epoch 72/80
5278/5278 [=====] - 1s 127us/step - loss: 1.2668e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 73/80
5278/5278 [=====] - 1s 117us/step - loss: 1.2537e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 74/80
5278/5278 [=====] - 1s 123us/step - loss: 1.2139e-07 - accuracy: 1.0000 - val_loss: 0.0013 - val_accuracy: 0.9992
Epoch 75/80
5278/5278 [=====] - 1s 122us/step - loss: 1.1772e-07 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 0.9992
Epoch 76/80
5278/5278 [=====] - 1s 111us/step - loss: 1.1609e-07 - accuracy: 1.0000 - val_loss: 0.0013 - val_accuracy: 0.9992
Epoch 77/80
5278/5278 [=====] - 1s 121us/step - loss: 1.1658e-07 - accuracy: 1.0000 - val_loss: 0.0011 - val_accuracy: 0.9992
Epoch 78/80
5278/5278 [=====] - 1s 119us/step - loss: 1.1436e-07 - accuracy: 1.0000 - val_loss: 0.0012 - val_accuracy: 0.9992
Epoch 79/80
5278/5278 [=====] - 1s 111us/step - loss: 1.1122e-07 - accuracy: 1.0000 - val_loss: 0.0011 - val_accuracy: 0.9992
Epoch 80/80
5278/5278 [=====] - 1s 119us/step - loss: 1.1408e-07 - accuracy: 1.0000 - val_loss: 9.7959e-04 - val_accuracy: 0.9992
dict_keys(['val_loss', 'val_accuracy', 'loss', 'accuracy'])
```

The bottom status bar shows the following information: Permissions: RW, End-of-lines: LF, Encoding: UTF-8, Line: 110, Column: 1, Memory: 68 %.

Performance:

- Accuracy: 0.9984

- F1 score :0.9950

```
epocn
Accuracy 0.9984848484848485
f1 score 0.9950248756218906


      precision    recall  f1-score   support

     0           1.00      1.00      1.00     1120
     1           0.99      1.00      1.00      200

 accuracy               1.00      1320
 macro avg              1.00      1320
weighted avg              1.00      1320
```

In [5]:

Accuracy graph:



Loss Graph:

